

Training cost. Per-sample gating cannot reduce the dense GPU cost —
a backward pass over a batch is the same price however many samples are masked.

